

practical_exercise_1_solution.Rmd

Lau Møller Andersen

6/9/2021

Solutions for practical exercise 1

Exercise 1

1. extract $\hat{\beta}$, Y , \hat{Y} , X and ϵ from **model** (hint: have a look at the function **model.matrix**)
 - i. create a plot that illustrates Y and \hat{Y} (if you are feeling ambitious, also include ϵ (hint:

```
data(mtcars)
model <- lm(mpg ~ wt, data=mtcars)

## extracting the parameters from the model object
print(beta.hat <- model$coefficients)

## (Intercept)          wt
##  37.285126    -5.344472

print(Y <- model$model$mpg)

## [1] 21.0 21.0 22.8 21.4 18.7 18.1 14.3 24.4 22.8 19.2 17.8 16.4 17.3 15.2 10.4
## [16] 10.4 14.7 32.4 30.4 33.9 21.5 15.5 15.2 13.3 19.2 27.3 26.0 30.4 15.8 19.7
## [31] 15.0 21.4

print(Y.hat <- model$fitted.values)

##           Mazda RX4      Mazda RX4 Wag      Datsun 710      Hornet 4 Drive
##           23.282611      21.919770      24.885952      20.102650
##   Hornet Sportabout      Valiant      Duster 360      Merc 240D
##           18.900144      18.793255      18.205363      20.236262
##           Merc 230      Merc 280      Merc 280C      Merc 450SE
##           20.450041      18.900144      18.900144      15.533127
##           Merc 450SL      Merc 450SLC  Cadillac Fleetwood  Lincoln Continental
##           17.350247      17.083024      9.226650      8.296712
##   Chrysler Imperial      Fiat 128      Honda Civic      Toyota Corolla
##           8.718926      25.527289      28.653805      27.478021
##           Toyota Corona  Dodge Challenger      AMC Javelin      Camaro Z28
##           24.111004      18.472586      18.926866      16.762355
##   Pontiac Firebird      Fiat X1-9      Porsche 914-2      Lotus Europa
##           16.735633      26.943574      25.847957      29.198941
##           Ford Pantera L      Ferrari Dino      Maserati Bora      Volvo 142E
##           20.343151      22.480940      18.205363      22.427495

print(X <- model.matrix(model))

##           (Intercept)      wt
## Mazda RX4           1  2.620
```

```
## Mazda RX4 Wag          1 2.875
## Datsun 710              1 2.320
## Hornet 4 Drive          1 3.215
## Hornet Sportabout      1 3.440
## Valiant                 1 3.460
## Duster 360             1 3.570
## Merc 240D              1 3.190
## Merc 230               1 3.150
## Merc 280               1 3.440
## Merc 280C              1 3.440
## Merc 450SE             1 4.070
## Merc 450SL             1 3.730
## Merc 450SLC            1 3.780
## Cadillac Fleetwood     1 5.250
## Lincoln Continental    1 5.424
## Chrysler Imperial      1 5.345
## Fiat 128               1 2.200
## Honda Civic            1 1.615
## Toyota Corolla         1 1.835
## Toyota Corona          1 2.465
## Dodge Challenger       1 3.520
## AMC Javelin            1 3.435
## Camaro Z28             1 3.840
## Pontiac Firebird       1 3.845
## Fiat X1-9              1 1.935
## Porsche 914-2          1 2.140
## Lotus Europa           1 1.513
## Ford Pantera L         1 3.170
## Ferrari Dino           1 2.770
## Maserati Bora          1 3.570
## Volvo 142E             1 2.780
## attr("assign")
## [1] 0 1
```

```
print(epsilon <- model$residuals)
```

```
##      Mazda RX4      Mazda RX4 Wag      Datsun 710      Hornet 4 Drive
##      -2.2826106      -0.9197704      -2.0859521      1.2973499
##      Hornet Sportabout      Valiant      Duster 360      Merc 240D
##      -0.2001440      -0.6932545      -3.9053627      4.1637381
##      Merc 230      Merc 280      Merc 280C      Merc 450SE
##      2.3499593      0.2998560      -1.1001440      0.8668731
##      Merc 450SL      Merc 450SLC      Cadillac Fleetwood      Lincoln Continental
##      -0.0502472      -1.8830236      1.1733496      2.1032876
##      Chrysler Imperial      Fiat 128      Honda Civic      Toyota Corolla
##      5.9810744      6.8727113      1.7461954      6.4219792
##      Toyota Corona      Dodge Challenger      AMC Javelin      Camaro Z28
##      -2.6110037      -2.9725862      -3.7268663      -3.4623553
##      Pontiac Firebird      Fiat X1-9      Porsche 914-2      Lotus Europa
##      2.4643670      0.3564263      0.1520430      1.2010593
##      Ford Pantera L      Ferrari Dino      Maserati Bora      Volvo 142E
##      -4.5431513      -2.7809399      -3.2053627      -1.0274952
```

```
## plotting Y and Y.hat together with the errors
par(font.lab=2, font.axis=2, cex=1.2)
```


```

plot(mtcars$wt, Y, ylim=c(0, 40), xlab='Weight (lb/1000)',
     ylab='Miles/(US) gallon', main='Plotting Y, Y.hat and epsilon')
lines(mtcars$wt, Y.hat)
n.obs <- dim(mtcars)[1]
for(index in 1:n.obs)
{
  x <- mtcars$wt[index]
  y0 <- Y[index]
  y1 <- Y.hat[index]

  lines(c(x, x), c(y0, y1), lty=3)
}

```

Plotting Y, Y.hat and epsilon



2. estimate β for a quadratic model ($y = \beta_2 x^2 + \beta_1 x + \beta_0$) using ordinary least squares *without* using **lm**;
 $\hat{\beta} = (X^T X)^{-1} X^T Y$ (hint: add a third column to X from step 1)

```

X.quad <- cbind(X, X[, 2]^2)
beta.hat.quad <- solve(t(X.quad) %*% X.quad) %*% t(X.quad) %*% Y
model.quad <- lm(mpg ~ wt + I(wt^2) + 1, data=mtcars)
print(beta.hat.quad)

```

```

##           [,1]
## (Intercept) 49.930811
## wt         -13.380337
##           1.171087

```

3. compare your acquired $\hat{\beta}$ with the output of the corresponding quadratic model created using **lm**

(hint: use the function **I**, see details under help and the sub-section formula operators here: <https://www.datacamp.com/community/tutorials/r-formula-tutorial>)

- i. create a plot that illustrates Y and \hat{Y} (if you are feeling ambitious, also include ϵ (hint: you can use the function **arrows**))

```
print(model.quad)


##
## Call:
## lm(formula = mpg ~ wt + I(wt^2) + 1, data = mtcars)
##
## Coefficients:
## (Intercept)          wt      I(wt^2)
##      49.931      -13.380       1.171

## plotting Y and Y.hat together with the errors
par(font.lab=2, font.axis=2, cex=1.2)
plot(mtcars$wt, Y, ylim=c(0, 40), xlab='Weight (lb/1000)',
      ylab='Miles/(US) gallon', main='Plotting Y, Y.hat and epsilon')
Y.hat <- model.quad$fitted.values
## we need to sort the values to use "lines"
sort.list <- sort(mtcars$wt, index.return=TRUE)
lines(sort.list$x, Y.hat[sort.list$ix])

n.obs <- dim(mtcars)[1]
for(index in 1:n.obs)
{
  x <- mtcars$wt[index]
  y0 <- Y[index]
  y1 <- Y.hat[index]

  lines(c(x, x), c(y0, y1), lty=3)
}
```

Plotting Y, Y.hat and epsilon



Exercise 2

1. which seems better?

It seems the fit is better for the quadratic fit

2. calculate the sum of squared errors, (show the calculation based on ϵ). Which fit has the lower sum?

```
epsilon.linear <- model$residuals
epsilon.quad <- model.quad$residuals

print(SS.linear <- sum(epsilon.linear^2))
```

```
## [1] 278.3219
```

```
print(SS.quad <- sum(epsilon.quad^2))
```

```
## [1] 203.7454
```

The sum of squared errors is smaller for the quadratic fit

3. now make a cubic fit ($y = \beta_3 x^3 + \beta_2 x^2 + \beta_1 x + \beta_0$) and compare it to the quadratic fit
 - i. create a plot that illustrates Y and \hat{Y} for both the cubic and the quadratic fits (plot them in the same plot)
 - ii. compare the sum of squared errors
 - iii. what's the estimated value of the "cubic" (β_3) parameter? Comment on this!

```


model.cub <- lm(mpg ~ I(wt^3) + I(wt^2) + wt + 1, data=mtcars)
## plotting Y and Y.hat together with the errors
par(font.lab=2, font.axis=2, cex=1.2)
plot(mtcars$wt, Y, ylim=c(0, 40), xlab='Weight (lb/1000)',
     ylab='Miles/(US) gallon', main='Plotting Y, Y.hat and epsilon')
Y.hat.quad <- model.quad$fitted.values
Y.hat.cub <- model.cub$fitted.values
## we need to sort the values to use "lines"
sort.list. <- sort(mtcars$wt, index.return=TRUE)

lines(sort.list$x, Y.hat.quad[sort.list$ix])
lines(sort.list$x, Y.hat.cub[sort.list$ix], lty=3, col='red')

legend('topright', c('Quadratic fit', 'Cubic fit'), col=c('black', 'red'),
      lwd=3)

```

Plotting Y, Y.hat and epsilon



```

epsilon.cub <- model.cub$residuals
print(SS.quad <- sum(epsilon.quad^2))

## [1] 203.7454
print(SS.cub <- sum(epsilon.cub^2))

## [1] 203.6699
print(beta.hat.3 <- model.cub$coefficients[2])

##      I(wt^3)

```

```
## 0.04593618
```

SS.cub is smaller than SS.quad but only very little so, indicating that the cubic part doesn't contribute a lot, as is also indicated by the small size of beta.hat.3. Note, that this isn't a reliable measure if the other beta.hats also change a lot

4. bonus question: which summary statistic is the fitted value (*Intercept* or β_0 in $y = \beta_0$) below identical to?

```
model.intercept <- lm(mpg ~ 1, data=mtcars)
print(mean(mtcars$mpg))
```

```
## [1] 20.09062
```

```
print(model.intercept$coefficients)
```

```
## (Intercept)
##      20.09062
```

Exercise 3

1. plot the fitted values for **logistic.model**:
 - i. what is the relation between the **linear.predictors** and the **fitted_values** of the **logistic.model**

```
logit <- function(x) log(x / (1 - x))
inv.logit <- function(x) exp(x) / (1 + exp(x))

logistic.model <- glm(am ~ wt + 1, data=mtcars, family='binomial')

print(logistic.model$fitted.values)
```

##	Mazda RX4	Mazda RX4 Wag	Datsun 710	Hornet 4 Drive
##	8.172115e-01	6.157283e-01	9.373069e-01	2.897304e-01
##	Hornet Sportabout	Valiant	Duster 360	Merc 240D
##	1.415972e-01	1.320944e-01	8.905706e-02	3.108616e-01
##	Merc 230	Merc 280	Merc 280C	Merc 450SE
##	3.463470e-01	1.415972e-01	1.415972e-01	1.290453e-02
##	Merc 450SL	Merc 450SLC	Cadillac Fleetwood	Lincoln Continental
##	4.884439e-02	4.030126e-02	1.132870e-04	5.625028e-05
##	Chrysler Imperial	Fiat 128	Honda Civic	Toyota Corolla
##	7.729920e-05	9.603663e-01	9.960953e-01	9.905887e-01
##	Toyota Corona	Dodge Challenger	AMC Javelin	Camaro Z28
##	8.929547e-01	1.067855e-01	1.440604e-01	3.193258e-02
##	Pontiac Firebird	Fiat X1-9	Porsche 914-2	Lotus Europa
##	3.131644e-02	9.859916e-01	9.686009e-01	9.974064e-01
##	Ford Pantera L	Ferrari Dino	Maserati Bora	Volvo 142E
##	3.283593e-01	7.097094e-01	8.905706e-02	7.013497e-01

```
print(inv.logit(logistic.model$linear.predictors))
```


##	Mazda RX4	Mazda RX4 Wag	Datsun 710	Hornet 4 Drive
##	8.172115e-01	6.157283e-01	9.373069e-01	2.897304e-01
##	Hornet Sportabout	Valiant	Duster 360	Merc 240D
##	1.415972e-01	1.320944e-01	8.905706e-02	3.108616e-01
##	Merc 230	Merc 280	Merc 280C	Merc 450SE
##	3.463470e-01	1.415972e-01	1.415972e-01	1.290453e-02
##	Merc 450SL	Merc 450SLC	Cadillac Fleetwood	Lincoln Continental
##	4.884439e-02	4.030126e-02	1.132870e-04	5.625028e-05

```
## Chrysler Imperial      Fiat 128      Honda Civic      Toyota Corolla
## 7.729920e-05      9.603663e-01      9.960953e-01      9.905887e-01
## Toyota Corona      Dodge Challenger      AMC Javelin      Camaro Z28
## 8.929547e-01      1.067855e-01      1.440604e-01      3.193258e-02
## Pontiac Firebird      Fiat X1-9      Porsche 914-2      Lotus Europa
## 3.131644e-02      9.859916e-01      9.686009e-01      9.974064e-01
## Ford Pantera L      Ferrari Dino      Maserati Bora      Volvo 142E
## 3.283593e-01      7.097094e-01      8.905706e-02      7.013497e-01
```

```
par(font.lab=2, font.axis=2, cex=1.2)
plot(mtcars$wt, logistic.model$fitted.values, xlab='Weight (lb/1000)',
     ylab='Probability of having manual transmission',
     main='Logistic regression - fitted values')
```

Probability of having manual transmission

Logistic regression – fitted values




The fitted values are the linear predictors with the inv.logit applied

2. plot the logistic function, you've estimated based on your $\hat{\beta}$, (not just the fitted values). Use an *xlim* of (0, 7)
 - i. what's the interpretation of the estimated $\hat{\beta}_0$ (the *Intercept*)
 - ii. calculate the estimated probability that the Pontiac Firebird has automatic transmission, given its weight
 - iii. bonus question - plot the logistic function and highlight all the cars where we guessed wrongly, if we used the following "quantizer" function:

$$transmission_{guess} = \begin{cases} 1(manual), & \text{if } PR(y = 1) \geq 0.5 \\ 0(automatic), & \text{otherwise} \end{cases} \quad (1)$$


```
# question i
wt <- seq(0, 7, 0.1)
p.am <- inv.logit(logistic.model$coefficients[1] +
  wt * logistic.model$coefficients[2])
par(font.lab=2, font.axis=2, cex=1.2)
plot(wt, p.am, type='l', lwd=3, xlab='Weight (lb/1000)',
  ylab='Probability of having manual transmission',
  main='Logistic regression - estimated function')
```



```
print(inv.logit(logistic.model$coefficients[1])) ## the probability of a car with weight 0 (!) having m
```

```
## (Intercept)
## 0.9999941
```

```
# question ii
pf.index <- which(rownames(mtcars) == 'Pontiac Firebird')
print(p.pf <- inv.logit(logistic.model$coefficients[1] +
  mtcars$wt[pf.index] * logistic.model$coefficients[2]))
```

```
## (Intercept)
## 0.03131644
```

```
# question iii
```

```
plot(wt, p.am, type='l', lwd=3, xlab='Weight (lb/1000)',
  ylab='Probability of having manual transmission',
```

```

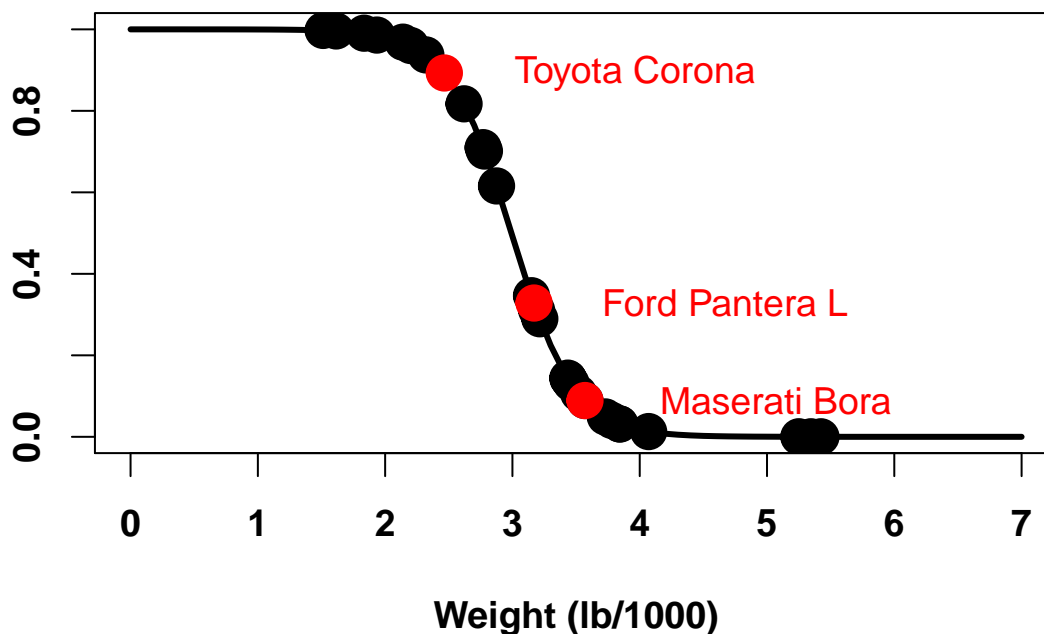
    main='Logistic regression, with wrong guesses highlighted')

mtcars$am.guess <- NA
n.obs <- dim(mtcars)[1]
for(index in 1:n.obs)
{
  temp.prob <- inv.logit(logistic.model$coefficients[1] +
    mtcars$wt[index] * logistic.model$coefficients[2])
  mtcars$am.guess[index] <- ifelse(temp.prob >= 0.5, 1, 0)
  colour <- ifelse(mtcars$am[index] == mtcars$am.guess[index],
    'black', 'red')
  if(colour == 'red') text(mtcars$wt[index] + 1.5, temp.prob,
    rownames(mtcars)[index], col=colour)
  points(mtcars$wt[index], temp.prob, col=colour, lwd=10)
}

```

Probability of having manual transmission

Logistic regression, with wrong guesses highlighted



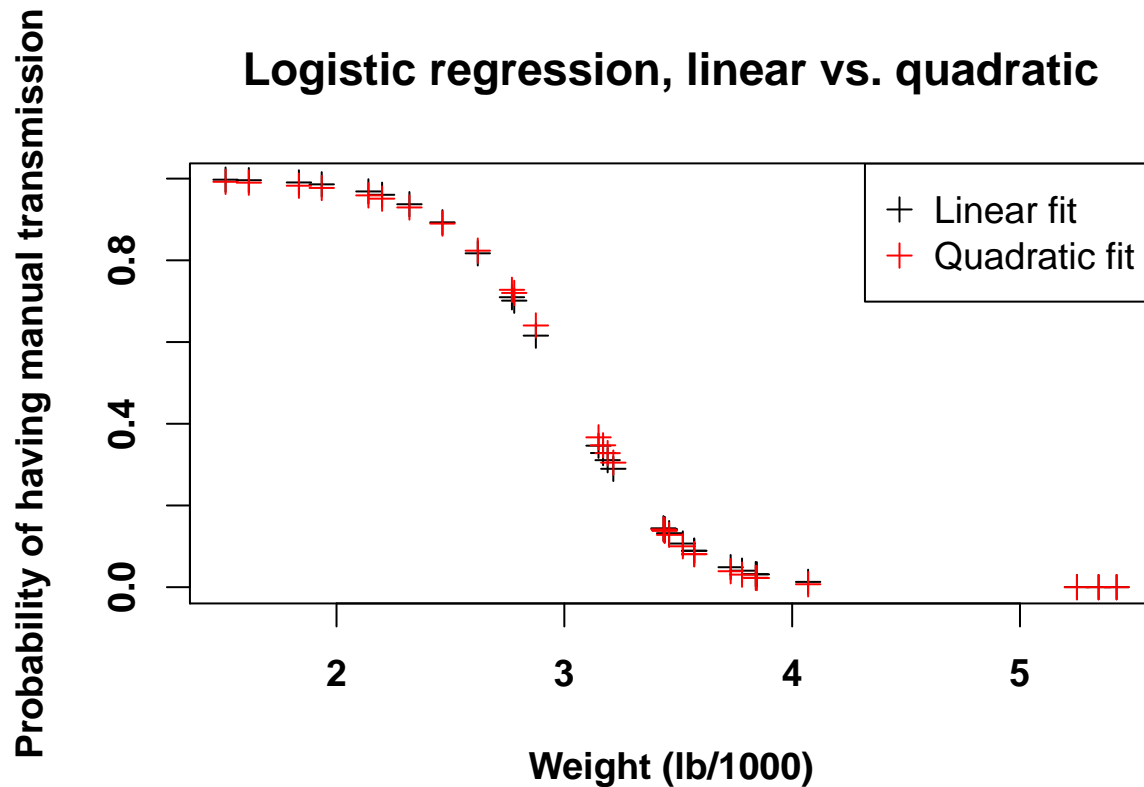
3. plot quadratic fit alongside linear fit
 - i. judging visually, does adding a quadratic term make a difference?
 - ii. check the details in the help of the AIC function - which of the models provide the better fit according to the AIC values and the residual deviance respectively?
 - iii. in your own words, why might it be good to penalise a model like the quadratic model, we just fitted.

```

quad.logistic.model <- glm(am ~ I(wt^2) + wt + 1, data=mtcars,
  family='binomial')
par(font.lab=2, font.axis=2, cex=1.2)

```

```
plot(mtcars$wt, logistic.model$fitted.values,
     xlab='Weight (lb/1000)',
     ylab='Probability of having manual transmission',
     main='Logistic regression, linear vs. quadratic', pch=3)
points(mtcars$wt, quad.logistic.model$fitted.values, col='red',
       pch=3)
legend('topright', c('Linear fit', 'Quadratic fit'), pch=3,
      col=c('black', 'red'))
```



```
# question i
# no, it doesn't make much of a difference
```

```
# question II
print(AIC(logistic.model))
```

```
## [1] 23.17608
```

```
print(AIC(quad.logistic.model))
```

```
## [1] 25.11779
```

```
# the linear model with only the linear term has the lower AIC, and is thus the better fit according to
```

```
# question iii
```

```
# adding more parameters always results in a better fit when looking at the residual deviance - we thus
```

```
# For example, fitting the perfect model
```

```


perfect.formula <- 'mpg ~ 1'
dim(mtcars)[1]

## [1] 32
for(index in 2:n.obs)
{
  perfect.formula <- paste(perfect.formula, ' + I(wt^',
                           index-1, ')', sep='')
}
perfect.formula <- as.formula(perfect.formula)
perfect.model <- lm(perfect.formula, data=mtcars)
sort.list. <- sort(mtcars$wt, index.return=TRUE)
plot(sort.list$x,
      perfect.model$fitted.values[sort.list$ix], type='l',
      ylim=c(0, 40), xlab='Weight (lb/1000)',
      ylab='Miles/(US) gallon', main='The "perfect" model')
points(mpg ~ wt, data=mtcars)
## we are not actually getting a perfect fit however...
print(SS.perfect <- sum(perfect.model$residuals^2))

## [1] 122.8369
# This is likely due to lack of numeric precision
library(Matrix)

```

The "perfect" model



```

## rank of matrix
print(rankMatrix(model.matrix(perfect.model))[1])

```

```
## [1] 7
```

```
## full rank would be 32
```