

Methods 3: Multilevel Statistical Modeling and Machine Learning

Week 5: *Evaluating and comparing models*
October 12, 2021

by: Lau Møller Andersen

These slides are distributed according to the CC
BY 4.0 licence:


<https://creativecommons.org/licenses/by/4.0/>



Study Cafe

Python book in Stakbogladen

<https://www.stakbogladen.dk/soegning.asp?phrase=9781783555130>



**I still owe the answers to questions
in the CryptPad
Poisson and overdispersion,
deviance in a logistic model, and the
assumption of normality

Code review – pair programming**

Learning goals

Evaluating and comparing models

1) Learning tools for comparing models

1) Variance explained

2) Likelihood ratio tests

3) Information criteria

2) Bridging to out-of-sample

Why are we modelling?

Remember Emil's slides from week 03

- To be able to understand the world
- To be able to predict and manipulate the world

gravitational pull equation
tells a truth about the world

$$F = G \frac{m_1 m_2}{r^2}$$

EXPLANATION

and allows us to --



NASA/Bill Ingalls

PREDICTION

What constitutes a good model?

Remember Emil's slides from week 03

- realist approach to science Accurate estimation of the underlying parameters of the population distribution
- Generalisation to new data we're not interested in truth, but in the ability to manipulating the world. If we use X tool, we should get Y effect...

EXPLANATION

PREDICTION

Within an **explanatory** framework, how can we assess whether we have done a good job?

1) Variance explained?

R summary

```
##
## Call:
## lm(formula = mpg ~ wt, data = mtcars)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -4.5432 -2.3647 -0.1252  1.4096  6.8727
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  37.2851     1.8776   19.858 < 2e-16 ***
## wt          -5.3445     0.5591   -9.559 1.29e-10 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 3.046 on 30 degrees of freedom
## Multiple R-squared:  0.7528, Adjusted R-squared:  0.7446
## F-statistic: 91.38 on 1 and 30 DF, p-value: 1.294e-10
```

Only include the stuff in assignments, that we discuss / are relevant for assignment

R-squared: measure of variance explained (from 0-1)

R^2 (coefficient of determination)

$$R^2 = 1 - \frac{\sum_{i=1}^n (y_i - \hat{y}_i)^2}{\sum_{i=1}^n (y_i - \mu_y)^2}$$

R^2 ; adjusted

relevant for multilevel models

$$R^2 = 1 - \frac{\sum_{i=1}^n (y_i - \hat{y}_i)^2}{\sum_{i=1}^n (y_i - \mu_y)^2}$$

$$\bar{R}^2 = 1 - (1 - R^2) \frac{n-1}{n-p-1}$$

n : number of observations

p : number of predictors beyond the constant term

```
lm(formula = mpg ~ wt, data = mtcars)
```

Sammenligning af forskellige modeller og hvor meget varians de hver især forklarer (r-squared)

Coefficients:

(Intercept)	wt
37.285	-5.344

"R-squared: 0.753"

"R-squared, adjusted: 0.745"

```
lm(formula = mpg ~ I(wt^2) + wt, data = mtcars)
```

Coefficients:

(Intercept)	I(wt^2)	wt
49.931	1.171	-13.380

"R-squared: 0.819"

"R-squared, adjusted: 0.807"

but how do we compare them in other ways? Is the increase significant e.g.?

```
lm(formula = mpg ~ I(wt^3) + I(wt^2) + wt, data = mtcars)
```

Coefficients:

(Intercept)	I(wt^3)	I(wt^2)	wt
48.40370	0.04594	0.68938	-11.82598

"R-squared: 0.819"

"R-squared, adjusted: 0.8"

```
lm(formula = mpg ~ I(wt^4) + I(wt^3) + I(wt^2) + wt, data = mtcars)
```

Coefficients:

(Intercept)	I(wt^4)	I(wt^3)	I(wt^2)	wt
14.4558	-0.3863	5.2004	-23.7018	36.6195

"R-squared: 0.822"

"R-squared, adjusted: 0.796"

R² – mixed models

how much does the group level effect explain of variance

$$R_{GLMM}(m)^2 = \frac{\sigma_f^2}{\sigma_f^2 + \sigma_\alpha^2 + \sigma_\epsilon^2}$$

how much does group and random effects explain of variance

$$R_{GLMM}(c)^2 = \frac{\sigma_f^2 + \sigma_\alpha^2}{\sigma_f^2 + \sigma_\alpha^2 + \sigma_\epsilon^2}$$

σ_f^2 : variance of the fixed effects

σ_α^2 : variance of the random effects

σ_ϵ^2 : unexplained variance

Variance explained

- Pros
 - R^2 is intuitive
- Cons
 - More complex models will always explain more variance
 - Hard to interpret in the case of collinearity
 - R^2 doesn't give us what we want

As opposed of confidence intervals, which tells us something about the population but not the sample.

$$y_i = \beta_0 + \beta_1 X_{1i} + \beta_2 X_{2i} + \varepsilon_i. \quad (1)$$



$$\hat{y}_i = 1.6 + 0.35X_{1i} + 0.62X_{2i}.$$



$$R^2 = 0.750$$

r-squared only based on this sample
meaning that it is not necessarily generalizable
r-squared doesn't validate the parameters



Tempting interpretation: the parameter estimates are likely to be true, because R^2 is large

if you want to generalize, you can say this r-squared would be my best guess as to what r-squared would be in another sample. But we don't know for sure

Correct interpretation:

“if one fits a model with the form of equation 1 in each new sample – *each time estimating new values of b_0 , b_1 , and b_2* – what will be the average proportional reduction in the sum of squared errors?”
(Yarkoni and Westfall 2017)

2) Likelihood ratios

Maximum likelihood estimation


$\sigma^2 = \text{variance}$

$E = \text{a constant}$

$P_i = \text{constant}$

probability density functions

$$\text{PDF} = (2\pi\sigma^2)^{-1/2} e^{-\frac{1}{2}\left(\frac{x-\mu}{\sigma}\right)^2}$$



$$\text{Conditional PDF}(y_i|X) = (2\pi\sigma^2)^{-1/2} e^{-\frac{1}{2}\left(\frac{y_i - x_i\beta}{\sigma_0}\right)^2}$$

What's the probability density function for each of our variables?

$$\text{Likelihood function: } L(\sigma^2, \epsilon) = (2\pi\sigma^2)^{-N/2} e^{-\left(\frac{1}{2\sigma^2} \sum_{i=1}^N (y_i - x_i\beta)^2\right)}$$

$$\text{Log-Likelihood function: } l(\sigma^2, \epsilon) = \log(L)$$

y : dependent variable

$X\beta$: linear predictor

$y - X\beta = \epsilon$: residuals

N : number of observations

Remember to ensure independence of parameters/observations

```
model <- lmer(Reaction ~ days_deprived + (days_deprived | Subject),  
             data=sleepstudy)  
model.ranint <- lmer(Reaction ~ days_deprived + (1 | Subject),  
                    data=sleepstudy)
```

which of these models makes the observations least unlikely
the closer to 0 the better

```
print(ll.m <- logLik(model))
```

```
## 'log Lik.' -702.0472 (df=6)
```

and this measurement is generalized to the population.
“The more parameters you include in the model the less unlikely your
observations will be to occur, since there are more parameters to explain
them.”

```
exp(ll.m)
```

```
## 'log Lik.' 1.272865e-305 (df=6)
```

```
print(ll.mr <- logLik(model.ranint))
```

```
## 'log Lik.' -715.01 (df=4)
```

```
exp(ll.mr)
```

```
## 'log Lik.' 2.986092e-311 (df=4)
```

Likelihood-ratio test

$$LR = -2(l(\theta_2) - l(\theta_1))$$

difference between the log likelihood times 2
(to get a positive number)

```
anova(model.ranint, model.ranslope.and.int)
```

```
## Data: sleepstudy
## Models:
## model.ranint: Reaction ~ days_deprived + (1 | Subject)
## model.ranslope.and.int: Reaction ~ days_deprived + (days_deprived | Subject)
##
```

	Df	AIC	BIC	logLik	deviance	Chisq	Chi	Df			
## model.ranint	4	1446.5	1458.4	-719.25	1438.5						
## model.ranslope.and.int	6	1425.2	1443.0	-706.58	1413.2	25.332		2			
##											
## model.ranint											
## model.ranslope.and.int											
##											
## ---											
## Signif. codes:	0	'***'	0.001	'**'	0.01	'*'	0.05	'.'	0.1	' '	1

for the likelihood test, we get a significance for the difference between the model performances (as opposed to r-squared).

the chi-squared is the difference between the log likelihood times 2

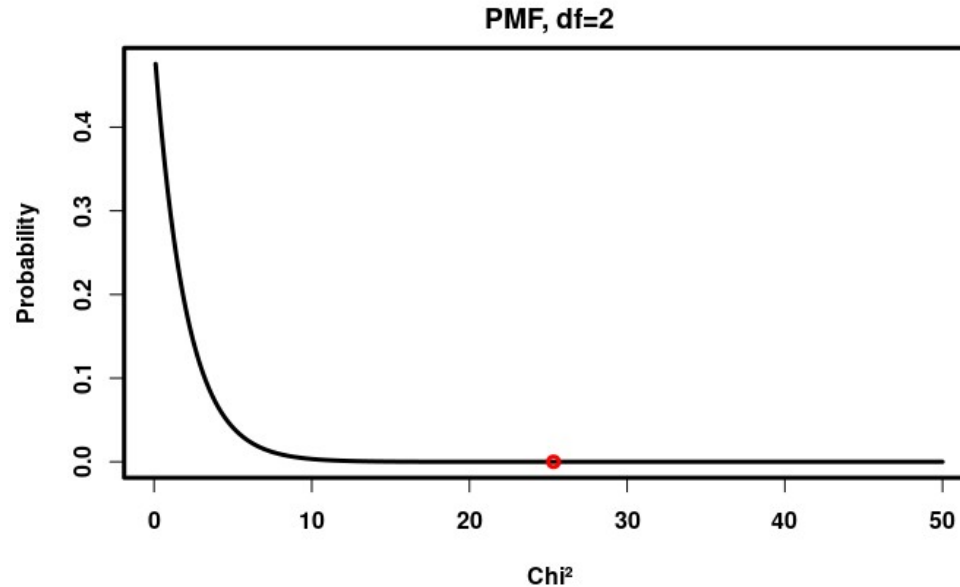
the p-value of the hypothesis to 'true' (thus enabling us to reject null hyp)

Likelihood-ratio test

$$LR = -2(l(\theta_2) - l(\theta_1))$$

Chisq	Chi	Df
25.332		2

Pr(>Chisq)
3.156e-06

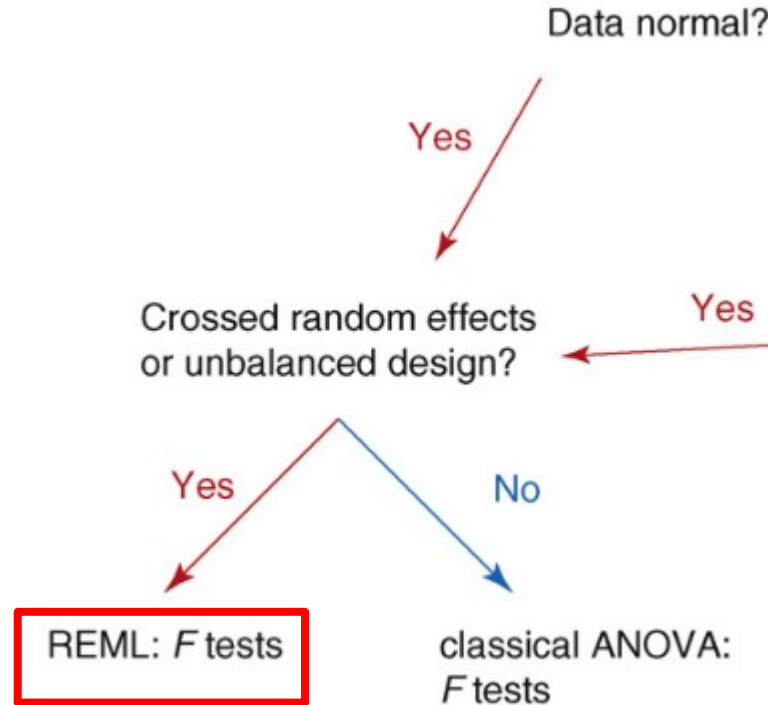


the higher chi-squared the more the model diverges from the null-hypothesis.

Regarding GLMMs

“The LR test is only adequate for testing fixed effects when both the ratio of the total sample size to the number of fixed-effect levels being tested and the number of random-effect levels (blocks) are large. *We have found little guidance and no concrete rules of thumb in the literature [...]*” (my highlights, Bolker et al., 2008)

**Requires
that
degrees
of
freedom
can be
defined**



(Bolker et al., 2009)

F-values – single level modelling

$$F = \frac{\left(\frac{RSS_1 - RSS_2}{p_2 - p_1} \right)}{\left(\frac{RSS_2}{n - p_2} \right)}$$

→ difference in unexplained variance between “null” model and “target” model

→ numerator degrees of freedom (difference in predictor variables)

→ unexplained variance by “target” model

→ denominator degrees of freedom (degrees of freedom for observations)

RSS : Residual sum of squares
 p : number of predictor variables
 n : number of observations

F-values – single level modelling


```
##  
## Call:  
## lm(formula = mpg ~ wt, data = mtcars) → Target model  
##  
## Residuals:  
##      Min       1Q   Median       3Q      Max   
## -4.5432 -2.3647 -0.1252  1.4096  6.8727   
##  
## Coefficients:  
##              Estimate Std. Error t value Pr(>|t|)      
## (Intercept)  37.2851     1.8776  19.858  < 2e-16 ***  
## wt          -5.3445     0.5591  -9.559 1.29e-10 ***  
## ---  
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1  
##  
## Residual standard error: 3.046 on 30 degrees of freedom  
## Multiple R-squared:  0.7528, Adjusted R-squared:  0.7446   
## F-statistic: 91.38 on 1 and 30 DF,  p-value: 1.294e-10
```

Call:
lm(formula = mpg ~ 1, data = mtcars)



Null model

F-values – single level modelling



```
## F-statistic: 91.38 on 1 and 30 DF,  p-value: 1.294e-10
```

F-values – single level modelling

```
##  
## Call:  
## lm(formula = mpg ~ I(wt^2) + wt, data = mtcars)  
##  
## Residuals:  
##      Min       1Q   Median       3Q      Max   
## -3.483  -1.998  -0.773   1.462   6.238   
##  
## Coefficients:  
##              Estimate Std. Error t value Pr(>|t|)      
## (Intercept)  49.9308     4.2113   11.856 1.21e-12 ***  
## I(wt^2)       1.1711     0.3594    3.258 0.00286 **   
## wt          -13.3803     2.5140   -5.322 1.04e-05 ***  
## ---  
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1  
##  
## Residual standard error: 2.651 on 29 degrees of freedom  
## Multiple R-squared:  0.8191, Adjusted R-squared:  0.8066  
## F-statistic: 65.64 on 2 and 29 DF, p-value: 1.715e-11
```




Target model

```
Call:  
lm(formula = mpg ~ 1, data = mtcars)
```



Null model,
(still!)

F-values – single level modelling



F-statistic: 65.64 on 2 and 29 DF, p-value: 1.715e-11

```
anova(model.1, model.2)
```

```
## Analysis of Variance Table
```

```
##
```

```
## Model 1: mpg ~ wt
```



Null model

```
## Model 2: mpg ~ I(wt^2) + wt
```



Target model

```
##   Res.Df    RSS Df Sum of Sq      F Pr(>F)
```

```
## 1      30 278.32
```

```
## 2      29 203.75  1    74.576 10.615 0.00286 **
```

```
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

What's the null hypothesis?

Likelihood ratio

- Pros
 - Models can be compared in a principled way by reference to a theoretical distribution, χ^2 . (In the single level case, F can be calculated)
- Cons
 - Models have to be nested in one another
 - Maximum likelihood fitting may be biased for complex models
 - Requires large sample sizes
 - Be careful if collinearity is high

3) Information criteria

Information criteria

$$\text{deviance} = -2l(\hat{\theta})$$

$$\text{AIC} = \text{deviance} + 2k$$

k : number of predictors

When we add k predictors that are pure noise, deviance is reduced by an amount corresponding to a χ^2 distribution with k degrees of freedom.

(Gelman and Hill, 2006)

“On average, a predictor needs to reduce the deviance by 2 in order to improve the fit to new data”

(Gelman and Hill, 2006, p. 525)

```
## Data: sleepstudy
## Models:
## model.ranint: Reaction ~ days_deprived + (1 | Subject)
## model.ranslope.and.int: Reaction ~ days_deprived + (days_deprived | Subject)
##
```

	Df	AIC	BIC	logLik	deviance	Chisq	Chi	Df
## model.ranint	4	1446.5	1458.4	-719.25	1438.5			
## model.ranslope.and.int	6	1425.2	1443.0	-706.58	1413.2	25.332		2

```
## Pr(>Chisq)
## model.ranint
## model.ranslope.and.int 3.156e-06 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

How are the degrees of freedom calculated?

Hints: **Likelihood function**: $L(\sigma^2, \epsilon)$

$$\hat{\sigma}^2 = \frac{1}{N} \sum_{i=1}^N (y_i - x_i \hat{\beta}_N)^2$$

How to estimate k ?

With J parameters

- Complete pooling
 - all J parameters collapsed into 1 parameter, ($k = 1$)
- No pooling
 - each of the J parameters is estimated, ($k = J$)
- Partial pooling
 - the number of effective parameters may be estimated by sampling and the *Deviance Information Criterion* can be estimated (beyond this course)
 - the number of parameters, loosely speaking, depend on whether the parameters are estimated mostly by the group average or the individual's own average ($k = ?$)

Information criteria

- Pros
 - Models can be compared even though one is not nested within the other (response data has to be the same though)
- Cons
 - Number of effective parameters not well defined for multilevel models
 - Maximum likelihood fitting may be biased for complex models

Did you learn? (it's not easy)

Evaluating and comparing models

1) Learning tools for comparing models

- 1) Variance explained

- 2) Likelihood ratio tests

- 3) Information criteria

2) Bridging to out-of-sample

Learning goals

Explanation and prediction

- 1) Understanding that fitting (explaining) often leads to overfitting
- 2) Learning methods to prevent overfitting by introducing *bias*
- 3) Understanding how the error can be decomposed into *bias* and *variance*

$$\hat{\beta} = (X^T X)^{-1} X^T Y$$

maximises the likelihood of

$$Y = X\beta + \epsilon$$

for which link function?

To fit is to overfit

(Yarkoni and Westfall, 2017)

Overfitting: fitting sample-specific noise,
which is thus not representative of the
population






Fig. 1. Training and test error produced by fitting either a linear regression (left) or a 10th-order polynomial regression (right) when the true relationship in the population (red line) is linear. In both cases, the test data (green) deviate more from the model's predictions (blue line) than the training data (blue). However, the flexibility of the 10th-order polynomial model facilitates much greater overfitting, resulting in lower training error but much higher test error than the linear model. MSE = mean squared error.

(Yarkoni and Westfall, 2017)


A sample of 10 linear or quadratic?




A sample of 10 linear or quadratic?




A sample of 10 linear or quadratic?



Residuals (linear)



Residuals (quadratic)



Quadratic:

$$ax^2 + bx + c$$

Linear:

$$bx + c$$


Estimates

$a = 0,6184; b = 3,0201$




$b = 3,020$

Now a sample of 100



$$b = 1,650$$



$$a = -0,03074 \approx 0; b = 1,650$$

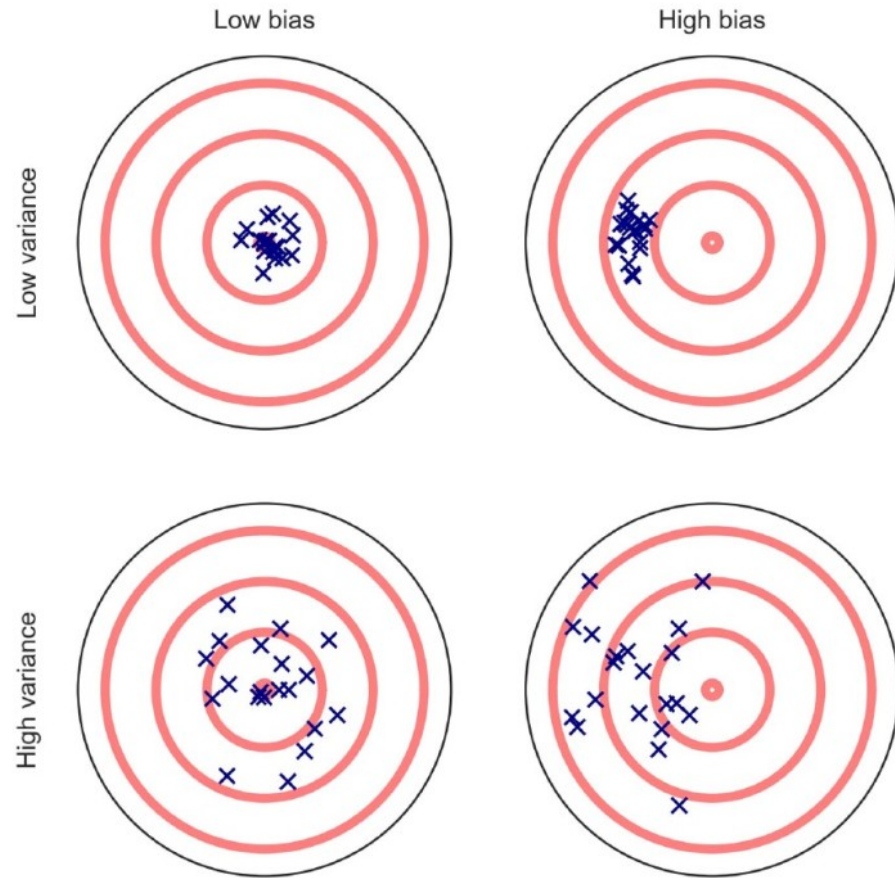


Fig. 2. An estimator's predictions can deviate from the desired outcome (or true scores) in two ways. First, the predictions may display a systematic tendency (or *bias*) to deviate from the central tendency of the true scores (compare right panels with left panels). Second, the predictions may show a high degree of *variance*, or imprecision (compare bottom panels with top panels).

(Yarkoni and Westfall, 2017)

Which two of these does our least-squares solution prefer?




FIGURE 3.7. Accuracy versus precision. A schematic illustration of the differences between accuracy and precision of source localization. After left median-nerve stimulation, activations is expected in the right-hemisphere hand region of the primary somatosensory cortex. The foot area is shown at the top of the head. See text for further explanation.

(Hari and Puce, 2017)




Fig. 3. Schematic illustration of the bias-variance decomposition. (Left) Under the classical error model, prediction error is defined as the sum of squared differences between true scores and observed scores (black lines). (Right) The bias-variance decomposition partitions the total sum of squared errors into two separate components: a bias term that captures a model's systematic tendency to deviate from the true scores in a predictable way (black line) and a variance term that represents the deviations of the individual observations from the model's expected prediction (gray lines).

Multilevel modelling as a *bias* introducer

“For example, some readers may be surprised to learn that multilevel modeling approaches to analyzing clustered data—which have recently seen a dramatic increase in adoption in psychology—improve on ordinary least squares (OLS) approaches to estimating individual cluster effects by deliberately biasing (through “shrinking” or “pooling”) the cluster estimates toward the estimated population average”

(Yarkoni and Westfall, 2017)

Introducing *bias*

“In a widely used form of penalized regression called lasso regression (Tibshirani, 1996, 2011), this leastsquares criterion is retained, but the overall cost function that the estimation seeks to minimize now includes an additional penalty term that is proportional to the sum of the absolute values of the coefficients.”

(Yarkoni and Westfall, 2017)

Penalised regression

$$\text{RSS} = \sum (y_i - \hat{y}_i)^2 \text{ (minimise to obtain least squares solution)}$$

$$\text{lasso regression: } \text{RSS} + \lambda \sum_{j=1}^p |\beta_j| \text{ (minimise this sum)}$$

$$\text{ridge regression: } \text{RSS} + \lambda \sum_{j=1}^p (\beta_j^2) \text{ (minimise this sum)}$$

i : observations

p : predictor variables

λ : a constant

Penalised regression

$RSS = \sum (y_i - \hat{y}_i)^2$ (minimise to obtain least squares solution)

lasso regression : $RSS + \lambda \sum_{j=1}^p |\beta_j|$ (minimise this sum)

ridge regression : $RSS + \lambda \sum_{j=1}^p (\beta_j^2)$ (minimise this sum)

i : observations

p : predictor variables

λ : a constant

Group discussion

In each case: what happens when?

1. λ increases?
2. λ decreases?
3. λ is 0?
4. λ goes towards infinity?

Penalised regression

$$\text{RSS} = \sum (y_i - \hat{y}_i)^2 \text{ (minimise to obtain least squares solution)}$$

$$\underset{\lambda}{\text{argmin}} = \text{RSS} + \lambda \sum_{j=1}^p |\beta_j|$$

$$\underset{\lambda}{\text{argmin}} = \text{RSS} + \lambda \sum_{j=1}^p (\beta_j^2)$$

i : observations

p : predictor variables

λ : a constant

How to choose λ ?

```
##  
## Call:  
## lm(formula = hp ~ mpg + wt + drat + qsec, data = mtcars)  
##  
## Coefficients:  
## (Intercept)          mpg          wt          drat          qsec  
##      473.779       -2.877       26.037        4.819       -20.751
```

What is λ equal to here?


```
sum.of.squares.total <- sum((y - mean(y))^2)  
sum.of.squared.errors.lm <- sum(residuals(linear_model)^2)  
print(r.squared.lm <- 1 - sum.of.squared.errors.lm/sum.of.squares.total)
```

```
## [1] 0.8072553
```

How to choose $\lambda_{(lasso)}$?

```
##  
## Call:  glmnet(x = x, y = y, alpha = 1, lambda = c(0, 0.2, 2, 4, 20,  
100))  
##  
##          lambda          RSS      penalty          sum  
## Df  %Dev Lambda          lambda          RSS      penalty          sum  
## 1  0 0.0000 100.0          100.0 145726.9          0.0          145726.9  
## 2  2 0.6567  20.0        20.00000 50025.64218        14.05496 50039.69714  
## 3  3 0.8004   4.0         4.00000 29082.92003        41.34363 29124.26366  
## 4  3 0.8051   2.0         2.00000 28408.50683        44.99057 28453.49740  
## 5  4 0.8072   0.2         0.20000 28097.60764        52.47741 28150.08505  
## 6  4 0.8073   0.0         0.00000 28088.09951        54.46997 28142.56948
```

How to choose λ ?



What does λ do?

$$\hat{\beta}_{ridge} = (X^T X + \lambda I)^{-1} X^T y$$

I : an identity matrix with p predictor variables

λ : a constant

What is $X^T X$?

```
head(X)
```

```
##                (Intercept)  mpg    wt drat  qsec
## Mazda RX4                1 21.0 2.620 3.90 16.46
## Mazda RX4 Wag            1 21.0 2.875 3.90 17.02
## Datsun 710                 1 22.8 2.320 3.85 18.61
## Hornet 4 Drive            1 21.4 3.215 3.08 19.44
## Hornet Sportabout         1 18.7 3.440 3.15 17.02
## Valiant                   1 18.1 3.460 2.76 20.22
```


```
print(cov.X)
```

$$X_{COV} = X^T X$$

```
##                mpg          wt          drat          qsec
## mpg  14042.310 1909.7528 2380.2770 11614.745
## wt   1909.753  360.9011  358.7190  1828.095
## drat 2380.277  358.7190  422.7907  2056.914
## qsec 11614.745 1828.0946 2056.9140 10293.480
```

$$X_{COV} = X^T X$$

Covariance matrix



The fact that the off-diagonal > 0 , indicates that there is **collinearity**


Collinearity can be bad

```
##  
## Call:  
## lm(formula = hp ~ mpg + wt + drat + qsec, data = mtcars)  
##  
## Coefficients:  
## (Intercept)          mpg           wt          drat          qsec  
##      473.779       -2.877       26.037        4.819       -20.751
```

Assuming no collinearity, what is the interpretation of the coefficients?
With collinearity, is that interpretation possible?


$$\hat{\beta}_{ridge} = (X^T X + \lambda I)^{-1} X^T y$$

Inverted Covariance matrix, regularized:(lambda= 0)




$$\hat{\beta}_{ridge} = (X^T X + \lambda I)^{-1} X^T y$$

Inverted Covariance matrix, regularized:(lambda= 1)




$$\hat{\beta}_{ridge} = (X^T X + \lambda I)^{-1} X^T y$$

Inverted Covariance matrix, regularized:(lambda= 10)




$$\hat{\beta}_{ridge} = \left(X^T X + \lambda I \right)^{-1} X^T y$$

Inverted Covariance matrix, regularized:(lambda= 100)



$$\hat{\beta}_{ridge} = (X^T X + \lambda I)^{-1} X^T y$$

Inverted Covariance matrix, regularized:(lambda= 1000)



**So why is it called
regularisation?**

Two notes about the inverted matrix:

with increase of λ

1. Diagonal shrinks (bias is added)
2. Off-diagonal shrinks (collinearity is reduced, which improves the stability of the model)

Stability of a model:

Feeding new data or adding new predictor variables will not change the parameter estimates a lot

We have succeeded in finding a λ making our model more stable (improved **in-sample** validity), but we haven't found a λ that optimises predictive power – (**out-of-sample**) – week 6)

Did you learn?

Explanation and prediction

- 1) Understanding that fitting (explaining) often leads to overfitting
- 2) Learning methods to prevent overfitting by introducing *bias*
- 3) Understanding how the error can be decomposed into *bias* and *variance*

References

- Bolker, B.M., Brooks, M.E., Clark, C.J., Geange, S.W., Poulsen, J.R., Stevens, M.H.H., White, J.-S.S., 2009. Generalized linear mixed models: a practical guide for ecology and evolution. *Trends in Ecology & Evolution* 24, 127–135.
<https://doi.org/10.1016/j.tree.2008.10.008>
- Gelman, A., Hill, J., 2006. *Data Analysis Using Regression and Multilevel/Hierarchical Models*. Cambridge University Press.
- Hari, R., Puce, A., 2017. *MEG-EEG Primer*. Oxford University Press, New York, NY, US.
- Yarkoni, T., Westfall, J., 2017. Choosing Prediction Over Explanation in Psychology: Lessons From Machine Learning. *Perspect Psychol Sci* 12, 1100–1122. <https://doi.org/10.1177/1745691617693393>